

Architecture principles, how to, patterns, examples

INTRODUCTION TO MICROSERVICES

October 2015, @SteveSfartz

AGENDA

- Why Microservices ?
- Designing Microservices
- Microservices Patterns
- Netflix stack
- Comparing SOA styles
- Sump up

Why, Monolithic style, Scaling the Monolith

MICROSERVICES ARCHITECTURE

Microservices (2014)

In short, the microservice architectural style is an approach to developing a **single application** as a **suite of small services**, each running in its own **process** and **communicating** with lightweight mechanisms, often an HTTP resource API.

These services are built around **business capabilities** and independently deployable by fully **automated deployment** machinery.

There is a bare minimum of centralized management of these services, which may be written in **different programming languages** and use different data storage technologies.

THE MONOLITHIC STYLE

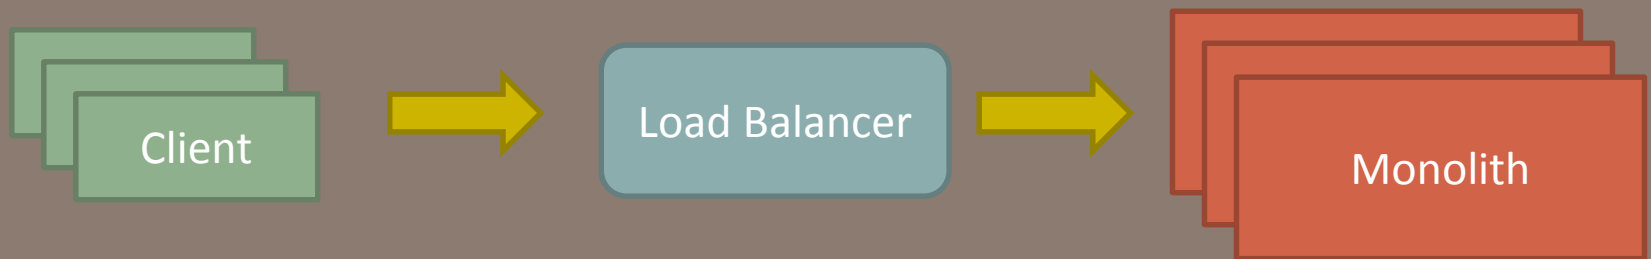
- An application built as a single unit, 3 main parts
 - a client-side user interface ,
 - a database
 - and a server-side application that handles HTTP requests, execute domain logic, retrieve and update data from the database, and select and populate HTML views to be sent to the browser.
- This server-side application is a *monolith*
 - a single **logical** executable
 - any change requires building and deploying a new version

THE MONOLITHIC STYLE

- Natural way to build a system
 - code specific to each layer (UI, Logic, DB)
 - generally a minimum of 3 languages
 - divided into classes, functions and namespaces
- Locally built and tested on devs' machines
- CI/CD pipeline to secure production

SCALING THE MONOLITH

- Several Monolith instances behind a load balancer

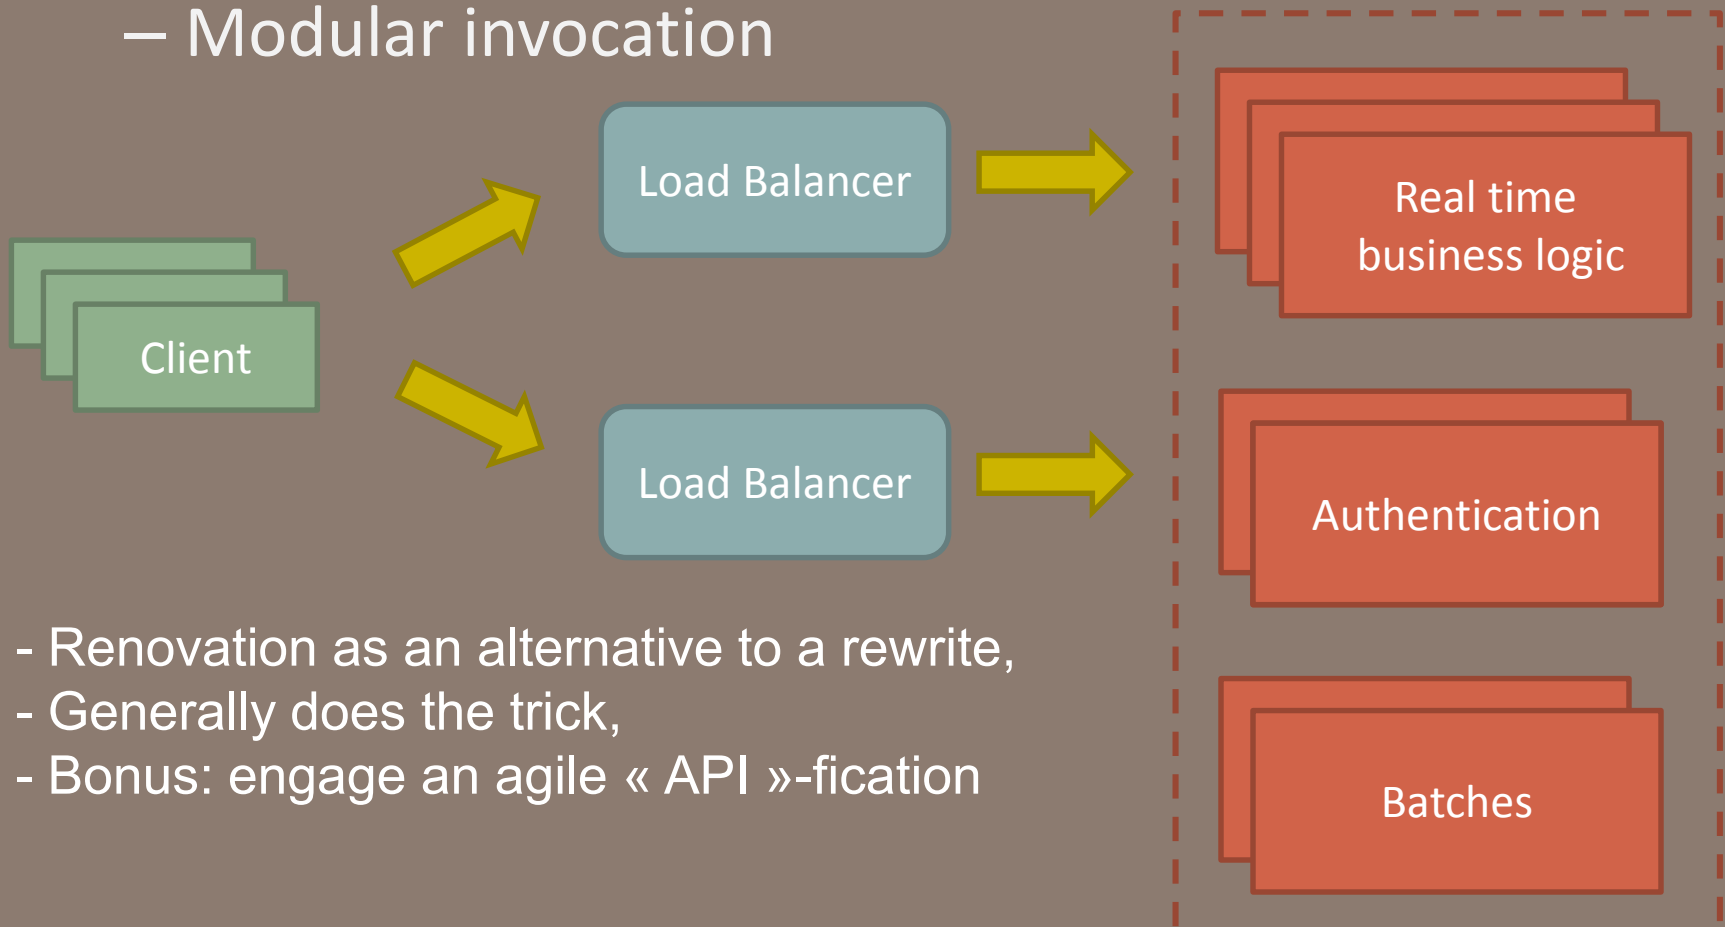


- + Quickest path to scale
- + High availability
- Routing traffic complexity
- Very large code base
- Change cycles tied together
- Limited scalability

⇒ lack of modularity

MAKING THE MONOLITH MORE MODULAR

- Specialized instances of a single codebase
 - Modular invocation



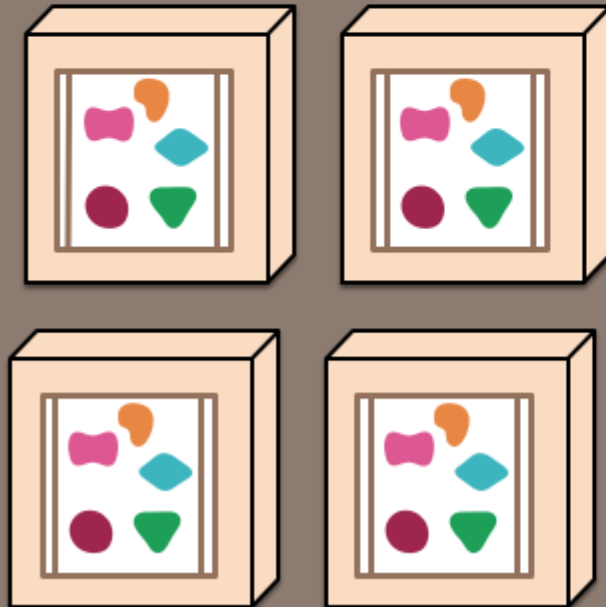
- Renovation as an alternative to a rewrite,
- Generally does the trick,
- Bonus: engage an agile « API »-fication

MONOLITHIC VS MICROSERVICES

A monolithic application puts all its functionality into a single process...



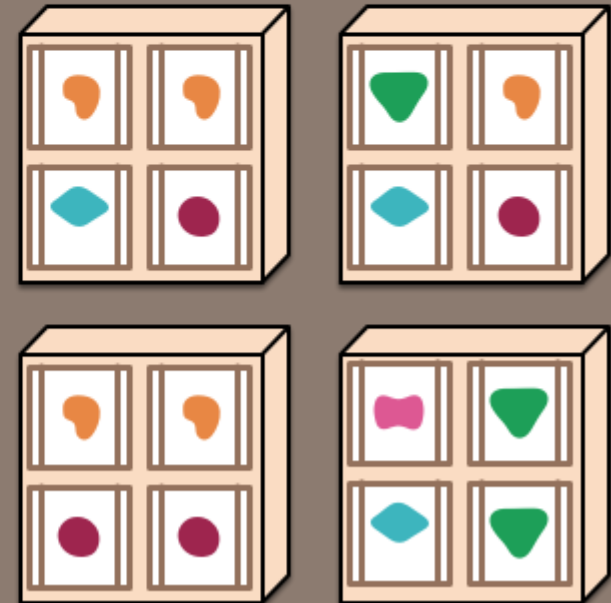
... and scales by replicating the monolith on multiple servers



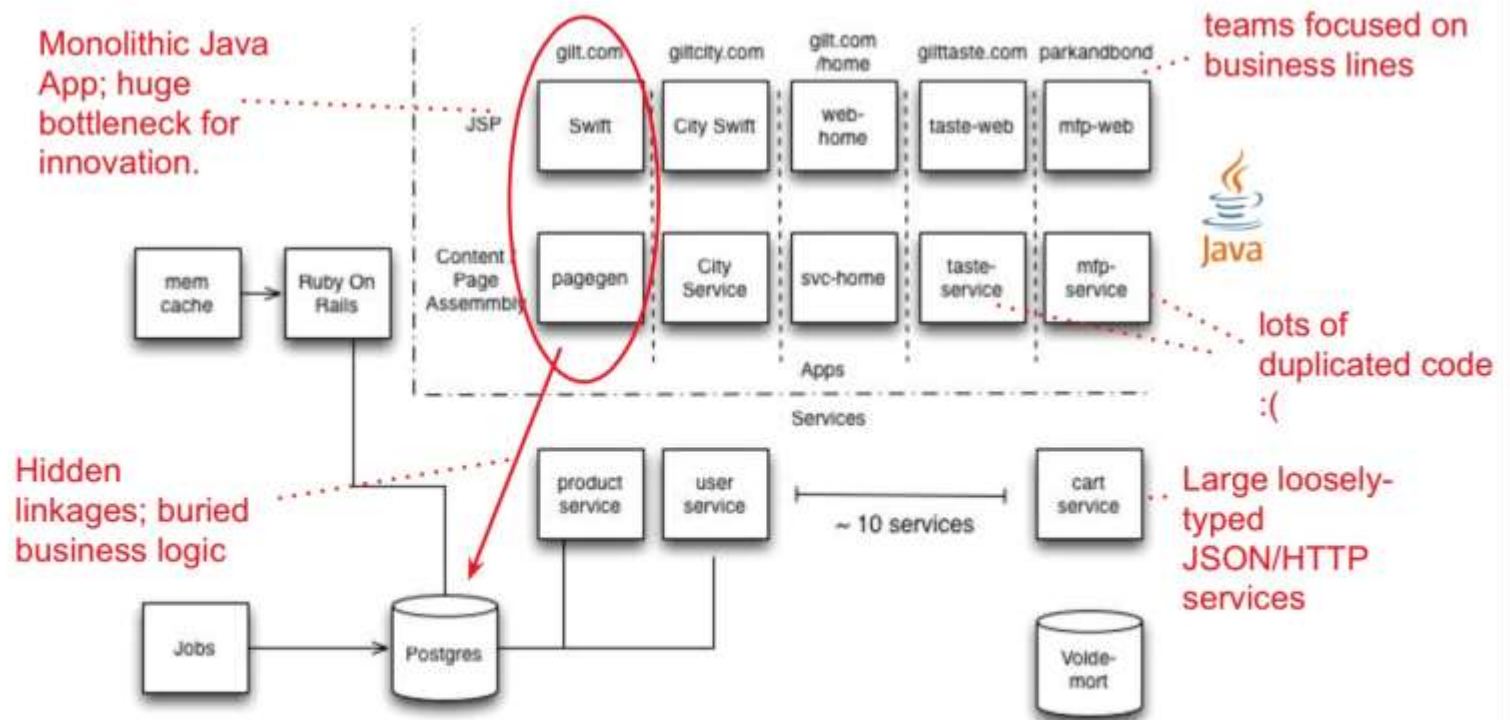
A microservices architecture puts each element of functionality into a separate service...



... and scales by distributing these services across servers, replicating as needed.

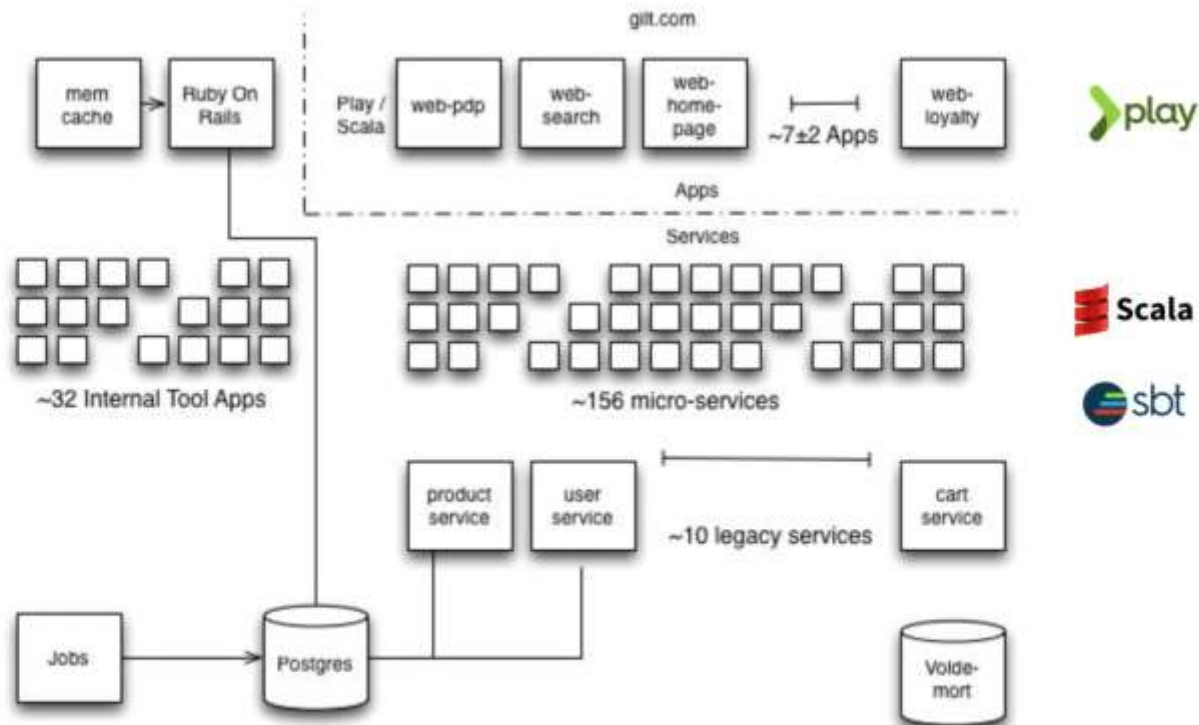


GILT TESTIMONIAL



2011: java, scala, loosely-typed services

GILT TESTIMONIAL



2015: micro-services

Common Characteristics

DESIGNING FOR MICROSERVICES

COMMON CHARACTERISTICS

- Componentization via Services
- Organized around Business Capabilities
- Products not Projects
- Smart endpoints and dumb pipes
- Decentralized governance
- Decentralized data management
- Infrastructure automation
- Design for failure
- Evolutionary design

COMPONENTIZATION VIA SERVICES

- Services as components rather than libraries
- Services avoid tight coupling by using explicit remote call mechanisms.
- Services are independently deployable and scalable

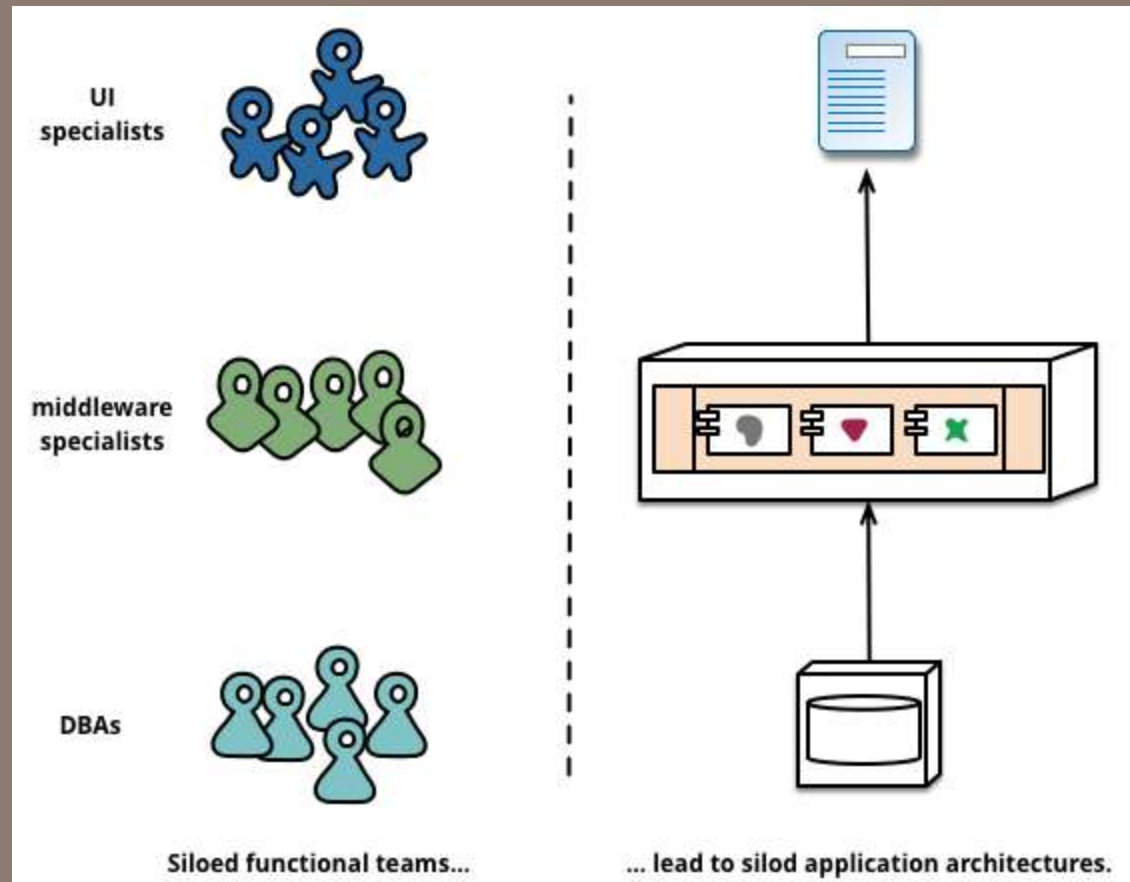
COMPONENTIZATION VIA SERVICES

- Each service also provides a firm module boundary
 - business or technical,
 - even allowing for different services to be written in different programming languages,
 - they can also be managed by different teams .
- A Microservice may consist of multiple processes
 - that will always be developed and deployed together,
 - Ex : an application process and a database that's only used by that service.

ORGANIZED AROUND BUSINESS CAPABILITIES


Any organization that designs a system will produce a design whose structure is a copy of the organization's communication structure.

Melvyn Conway, 1967.



ORGANIZED AROUND BUSINESS CAPABILITIES

- Microservices to solve Conway's anti-pattern



Cross-functional teams....

... organized around capabilities

PRODUCTS NOT PROJECTS

- Standard project model:
 - deliver pieces of software which are then considered to be completed,
 - hand over to a maintenance organization and disband the project team
- The Microservices style
 - a team should own a product over its full lifetime
 - Amazon : You build => You run it

SMART ENDPOINTS AND DUMB PIPES

- Be as decoupled and as cohesive as possible
 - own domain logic,
 - act more as filters in the classical Unix sense
 - using simple RESTish protocols and lightweight messaging
- Smarts live in the services, not in-between the endpoints
 - No central tool / bus that includes sophisticated routing, transformations, process, business rules
- Pre-requisite : turn the chatty in-process communication of the monolith into coarser-grained network messaging


DECENTRALIZED GOVERNANCE

- The common microservices practice is to choose the best tool for each job
 - Tools are shared inside and outside the organization
 - Focus on common problems : data storage, IPC, infrastructure automation
 - Example : Netflix opensource libraries
 - Favor independent evolution
 - Consumer-driven Service contracts,
 - Tolerant Reader Pattern
- ⇒ Minimal over-heads create an opportunity to have teams responsible for all aspects
- Build their microservices and operate with 24/7 SLAs

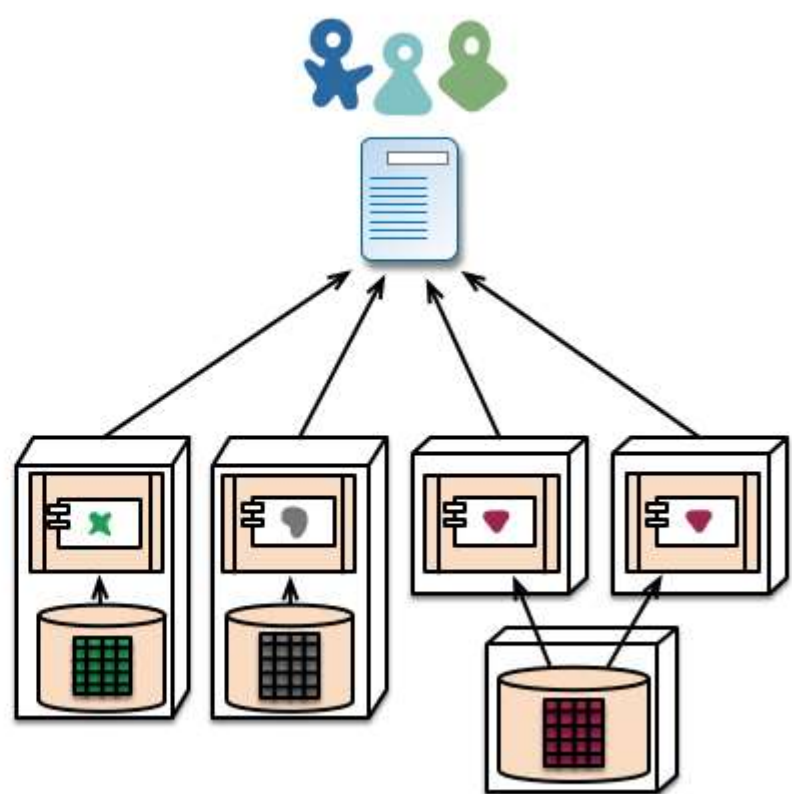
DECENTRALIZED DATA MANAGEMENT

- No unique data model approach
 - the conceptual model differs between microservices
 - Reinforce the separation of concerns
- Decentralized data storage
 - Polyglot persistence is frequent in microservices architectures
- Manage inconsistencies via the business practices in place throughout the organization
 - Common design : reversal processes versus 2PC distributed transactions

DECENTRALIZED DATA MANAGEMENT



monolith - single database



microservices - application databases

INFRASTRUCTURE AUTOMATION

- CI/CD, Real-time monitoring, global system and fine grained resources dashboards,
- Has become a standard practice
 - thanks to public cloud providers,
 - but also opensource tools
- a QA stake for monoliths
- a pre-requisite for microservices

DESIGN FOR FAILURE

- Any service call can fail
 - Circuit Breaker pattern,
 - async consumption
 - Best practice : 1 to ZERO sync call
 - Ex : Netflix Stack
 - Hystrix / Eureka / Ribbon
 - /!\ provide interop & implementations for various languages
- Infrastructure pre-requisites
 - monitor and restore at scale
 - dedicated tooling for simulation
 - Ex : Netflix Simian Army (Chaos Monkey)

EVOLUTIONARY DESIGN

- More granular release planning
- Goal : change tolerance
 - Identify change impacts quickly (via automation, service contracts, automated dependency maps)
 - Fix rather than revert
 - In case of mandatory breaking change, use versioning
 - a last resort option in the microservices world


A bit of history, Best practices, Netflix Stack

REAL WORLD MICROSERVICES

GOOGLE TRENDS – OCTOBER 2015

Interest over time ?

☒ News headlines ☐ forecast ?



Regional interest ?



Region | Town/City

India	100	<div><div></div></div>
United Kingdom	80	<div><div></div></div>
Germany	69	<div><div></div></div>
United States	64	<div><div></div></div>

STATE OF THE ART

- Pioneers : Amazon, Netflix, Google
 - Started Opensourcing their building blocks in 2010
 - Democratized the Microservices architecture style
- Today
 - Large corporation for complex systems that need to evolve frequently (understand continuously)
 - Digital startups to ensure their core business will scale and embrace the long tail

AMAZON DNA

- Microservices by design

1. All teams will henceforth expose their data and functionality through **service interfaces**.
2. Teams must communicate with each other through these interfaces.
3. There will be no other form of interprocess communication allowed: no direct linking, no direct reads of another team's data store, no shared-memory model, no back-doors whatsoever. The only communication allowed is via service interface calls over the network.
4. It doesn't matter what technology they use. HTTP, Corba, Pubsub, custom protocols — doesn't matter.
5. All service interfaces, without exception, must be designed from the ground up to be externalizable. That is to say, the team must plan and design to be able to **expose** the interface to developers in the outside world. No exceptions.
6. Anyone who doesn't do this will be fired.
7. Thank you; have a nice day!

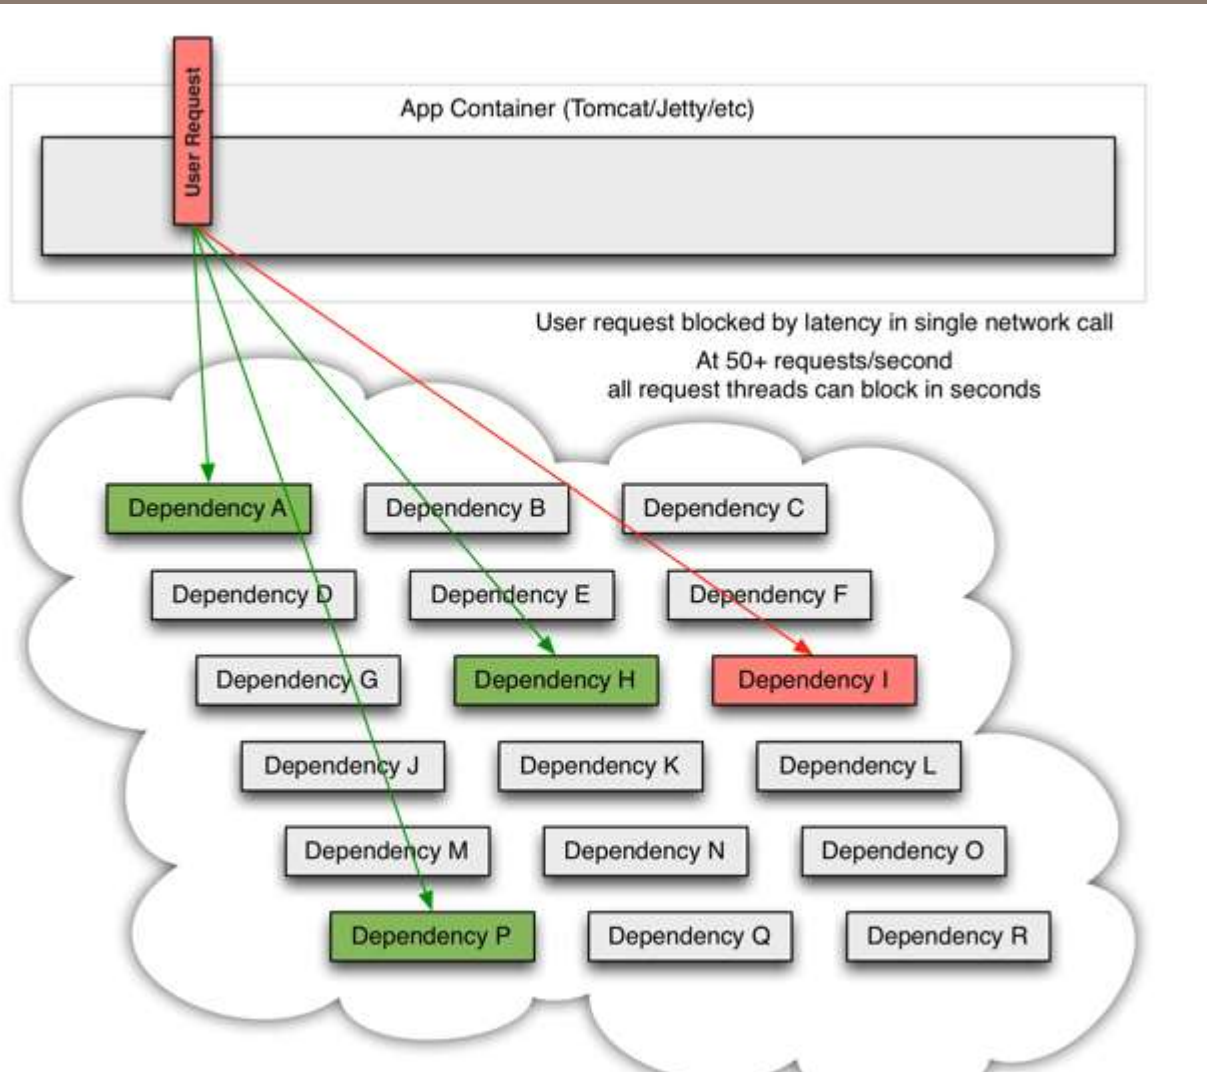
NETFLIX ECOSYSTEM

- 100s of microservices
- 1,000s of daily production changes
- 10,000s of instances
- 100,000s of customer interactions per minute
- 1,000,000s of customers
- 1,000,000,000s of metrics
- 10,000,000,000 hours of streamed
- 10s of operations engineers

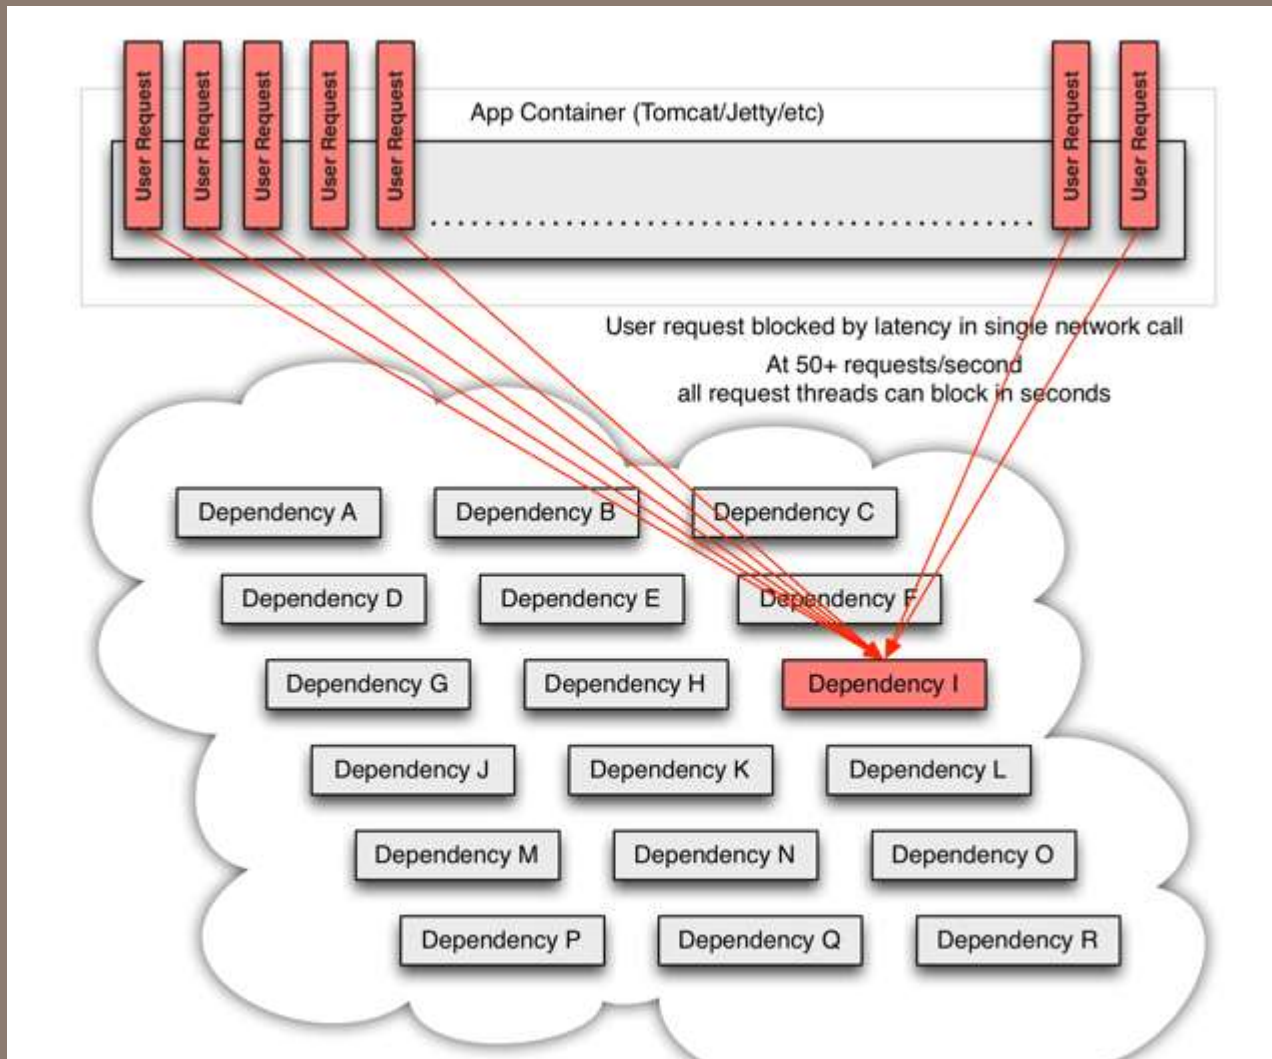
RESILIENCE IN DISTRIBUTED SYSTEMS

- Failures are common
- An application that depends on 30 services
 - where each service has 99.99% uptime
 - $99.99^{30} = 99.7\%$ uptime
 - 0.3% of 1 billion requests = 3,000,000 failures
 - 2+ hours downtime/month even if all dependencies have excellent uptime.
- Reality is generally worse.

WHEN 1 BACKEND BECOMES LATENT



AT 50+ REQ/S, ALL REQUEST THREADS CAN BLOCK IN SECONDS



FAIL FAST, FAIL SILENT, FALLBACK

- Wrapping all calls to external systems (or “dependencies”) which typically executes within a separate thread
- Timing-out calls that take longer than thresholds you define.
- Maintaining a small thread-pool for each dependency;
 - if it becomes full, requests immediately rejected instead of queued up.
- Measuring successes, failures (exceptions thrown by client), timeouts, and thread rejections.
- Tripping a circuit-breaker to stop all requests to a particular service for a period of time,
 - either manually or automatically (error percentage threshold)
- Performing fallback logic when a request fails, is rejected, times-out, or short-circuits.
- Monitoring metrics and configuration changes in near real-time.

SERVICES DISCOVERY

- Typically, a Microservice gets created and destroyed often
- It is reconfigured on the fly
- In near real-time, others should find it
- Examples
 - Apache ZooKeeper
 - HashiCorp Consul
 - CoreOS Etcd
 - Netflix Eureka

SERVICES CONSUMPTION

- Choose the best messaging infrastructure
 - RPC / REST style,
 - Request/Response, Streaming
 - HTTP, HTTP2, TCP, UDP
- Support async consumption / aggregation
 - Use parallel code structures of your clients
 - GO, ES7 (BabelJS, Google Traceur)
- Ease consumption via clients SDKs
 - Automated generated from API definitions
 - Service providers tend to opensource their SDKs
 - Facebook Parse, IBM Bluemix Mobile services

SERVICES CONSUMPTION

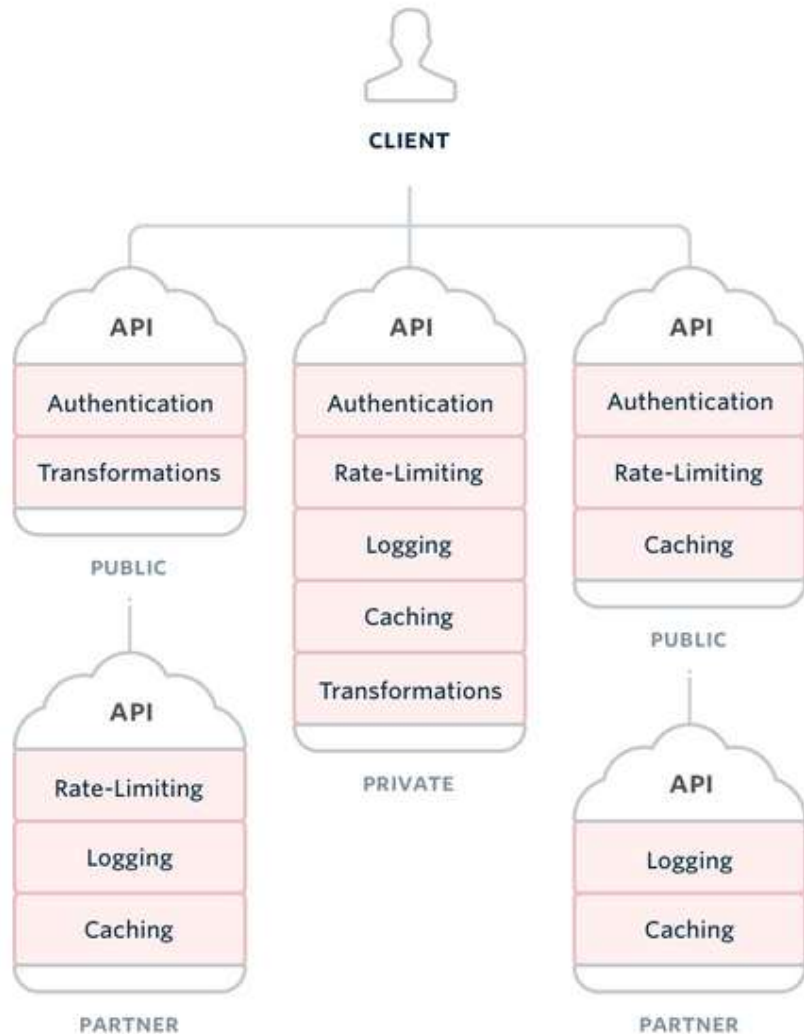
- Favor SDKs with large feature sets
 - Basic REST is not an option for microservices
 - integrate load balancing, fault tolerance, caching, service discovery, multiple transport protocols (HTTP, TCP, UDP)
- Netflix Ribbon
 - Inter Process Communication with built in software load balancers
 - Integrates Hystrix and Eureka
 - The primary usage model involves REST calls with various serialization scheme support.

MICROSERVICES GATEWAY

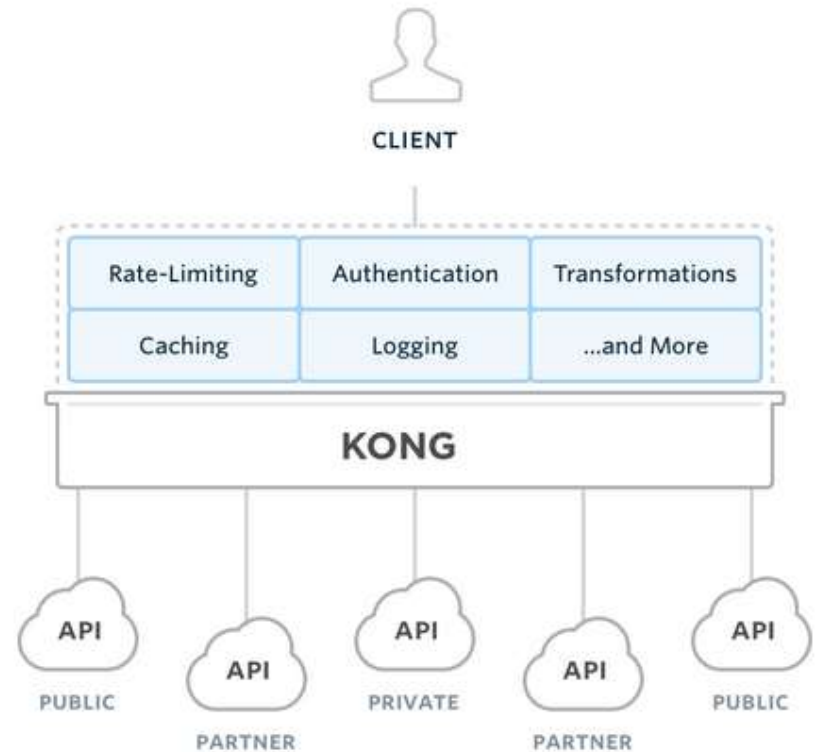
- As an Edge service to access back-end microservices
 - Authentication and Security
 - Dynamic routing,
 - Insights and Monitoring
 - Stress testing : gradually increasing the traffic to a cluster in order to gauge performance.
 - Load Shedding : allocating capacity for each type of request and dropping requests that go over the limit.
 - Static Response handling : building some responses directly at the edge instead of forwarding them to an internal cluster
 - Multiregion Resiliency : routing requests across AWS regions in order to diversify the ELB usage and move the edge closer to users

MICROSERVICES GATEWAY

- Lots of building blocks
 - Nginx
 - Mashape Kong
- Needs integration with the other building blocks of the Microservices communication infrastructure
 - Example : Netflix Zuul



- ✗ Common functionality is duplicated across multiple services
- ✗ Systems tend to be monolithic and hard to maintain
- ✗ Difficult to expand without impacting other services
- ✗ Productivity is inefficient because of system constraints



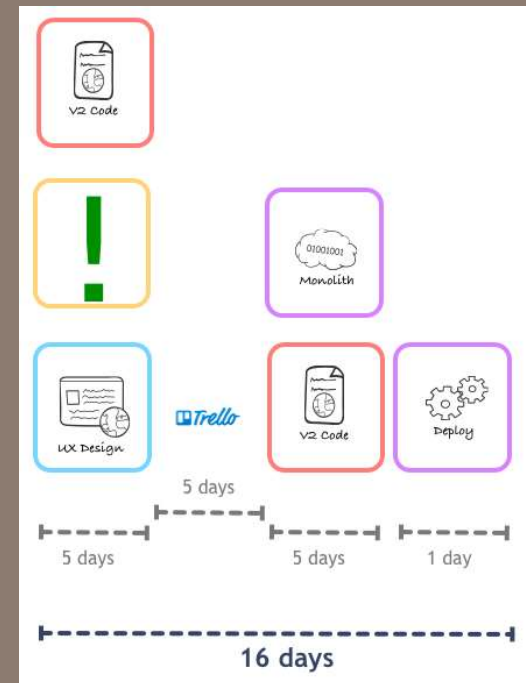
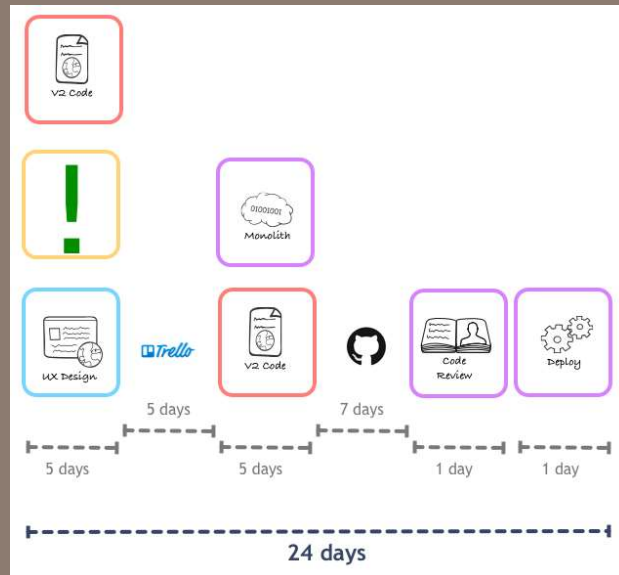
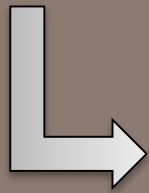
MASHAPE KONG

- ✓ Kong centralizes and unifies functionality into one place
- ✓ Build efficient distributed architectures ready to scale
- ✓ Expand functionality from one place with a simple command
- ✓ Your team is focused on the product, Kong does the REST

COMMITTED TEAMS

- Ownership core to team organization
 - built into the management of the organization
 - make sure that teams have sufficient time to truly own the applications that they are in charge
 - “Products versus Projects” principle
 - “Functional versus Divisional” organizations
 - And give them 360 view on operations

SOUNDCLOUD TESTIMONIAL



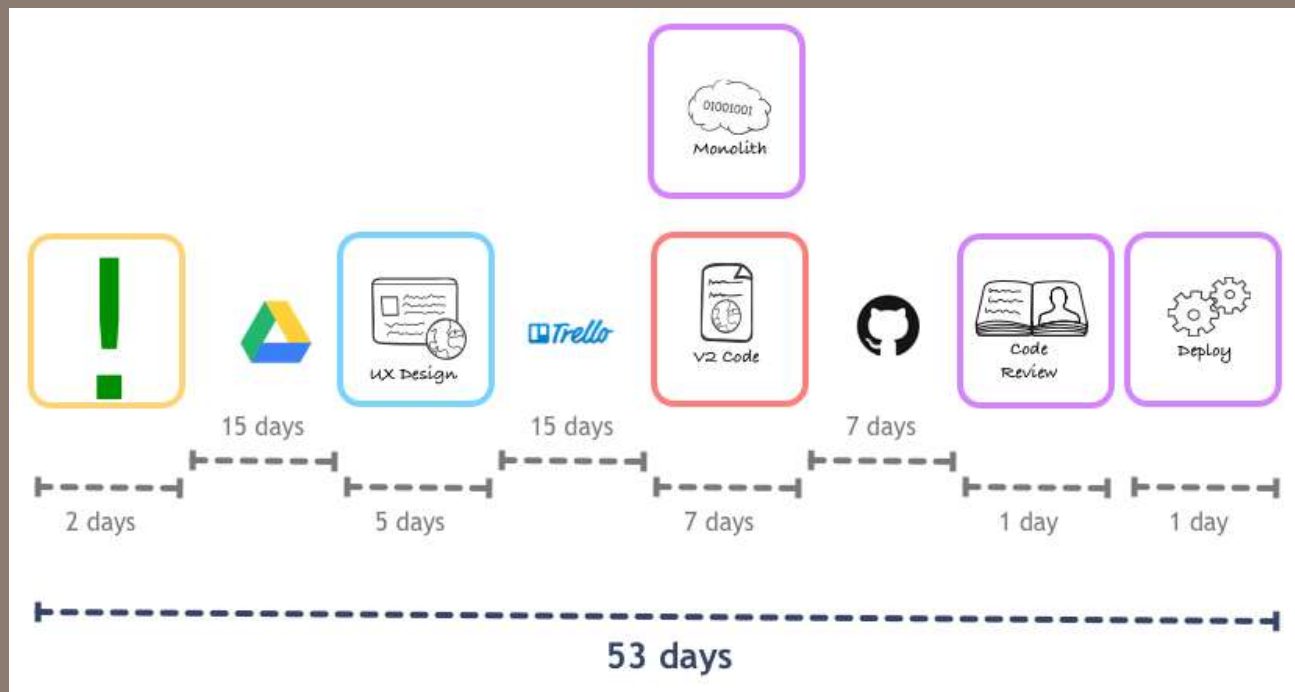
SOUNDCLOUD TESTIMONIAL

- Actual workflow to go live
 - 2 months, 11 steps
 - Main elephant is all the dance between front-end and back-end development
 - 11 days are doing actual development work



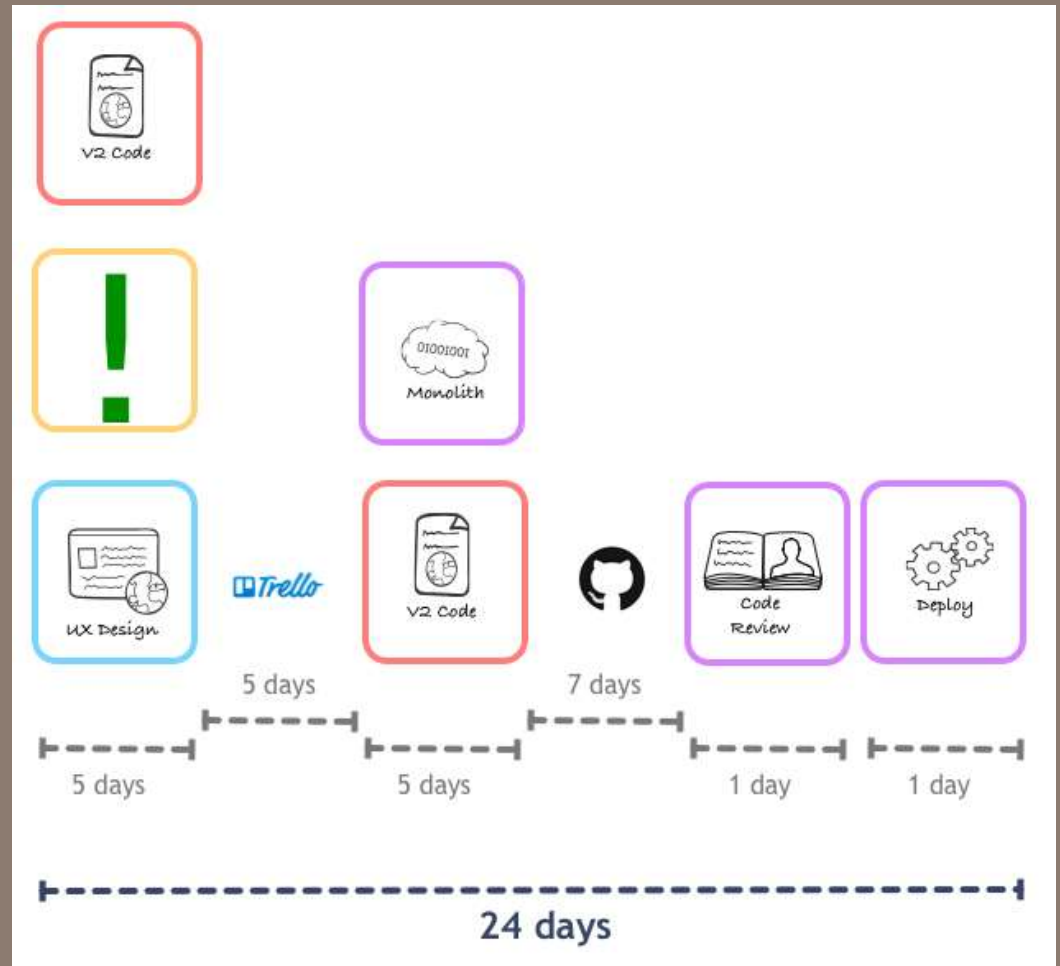
SOUNDCLOUD TESTIMONIAL

- Decision : pairing back-end and front-end devs
 - pair fully dedicated to a feature until its completion
 - Individually, each person ended up spending more time doing work per feature



SOUNDCLOUD TESTIMONIAL

- Designer, product manager, and front-end developer working close to each other

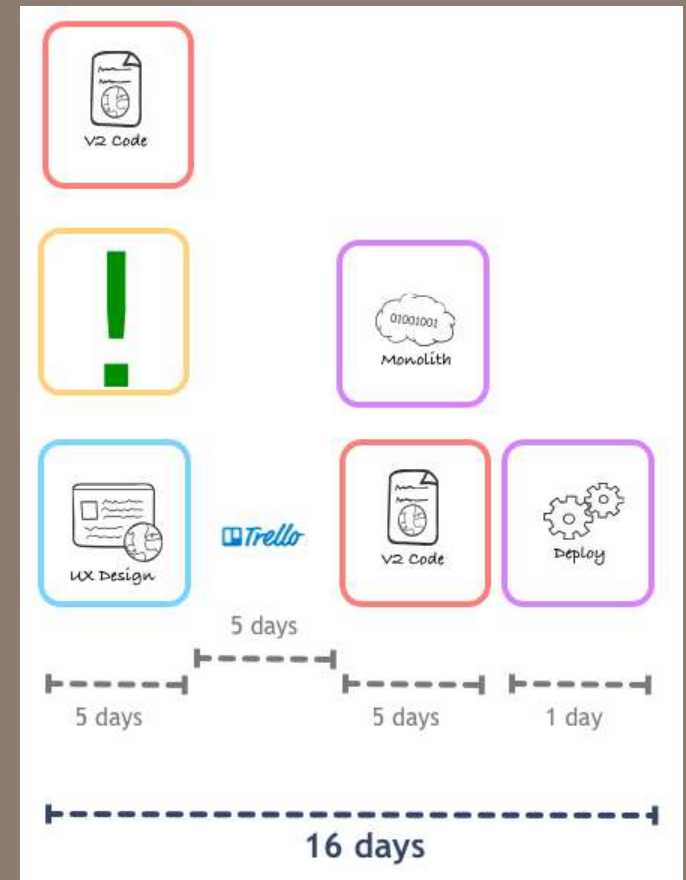


SOUNDCLOUD TESTIMONIAL

- The irreducible complexity of the monolith
 - *Why do we need Pull Requests?*
 - *Why do people make mistakes so often?*
 - *Why is the code base so complex?*
 - *Why do we need a single code base to implement the many components?*
 - *Why can't we have economies of scale for multiple, smaller, systems?*

SOUNDCloud TESTIMONIAL

- Isolated new features in dedicated microservices, isolated from the monolith
- New organization with team of 3 to 4 people
- Each team is responsible for decided whether parts of the Monolith are extracted and rewritten , or kept

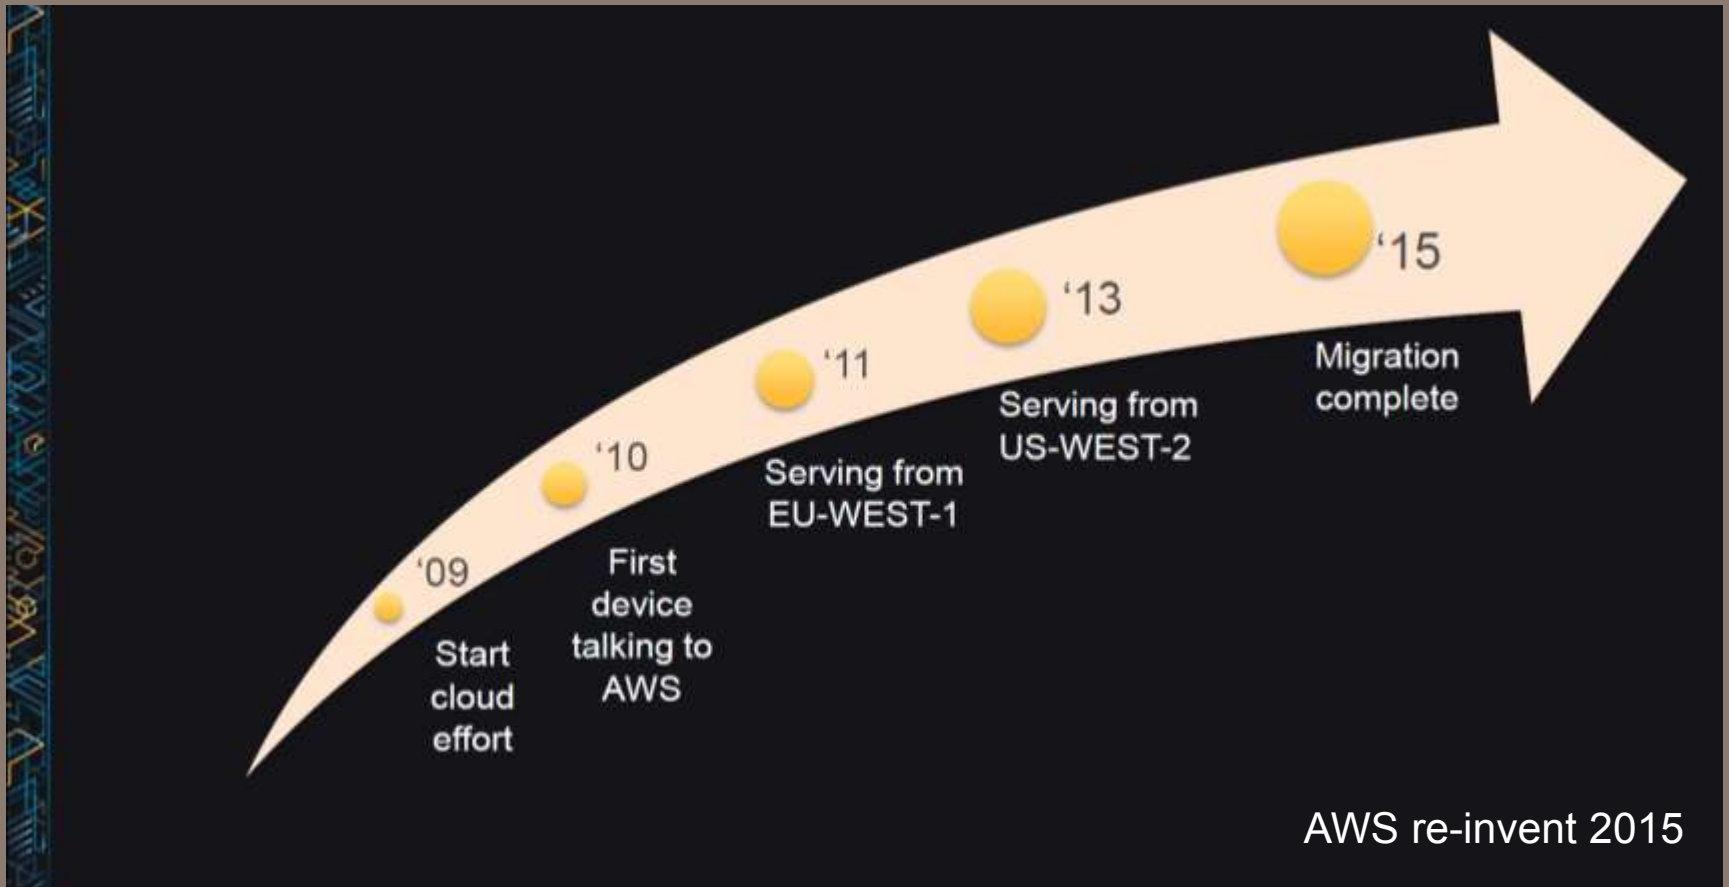


COMMITTED TEAMS

- Ownership core to team organization
 - built into the management of the organization
 - make sure that teams have sufficient time to truly own the applications that they are in charge
 - “Products versus Projects” principle
 - “Functional versus Divisional” organizations
 - And give them 360 view on operations

THE NETFLIX MICROSERVICES JOURNEY

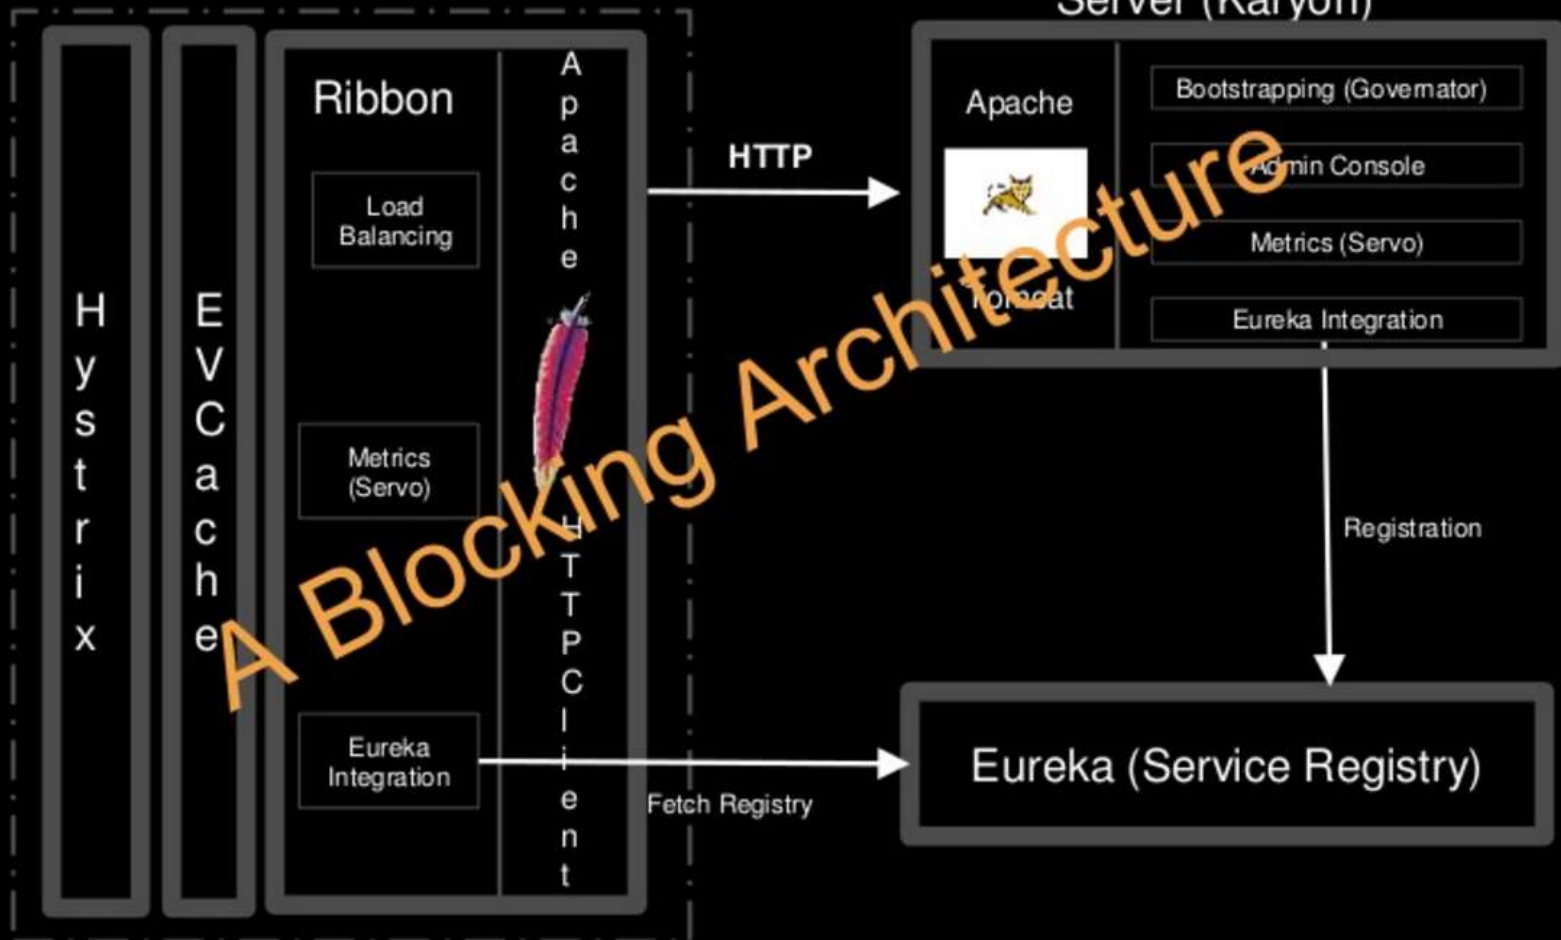
- Migration to AWS



Netflix IPC Stack (1.0)

Client

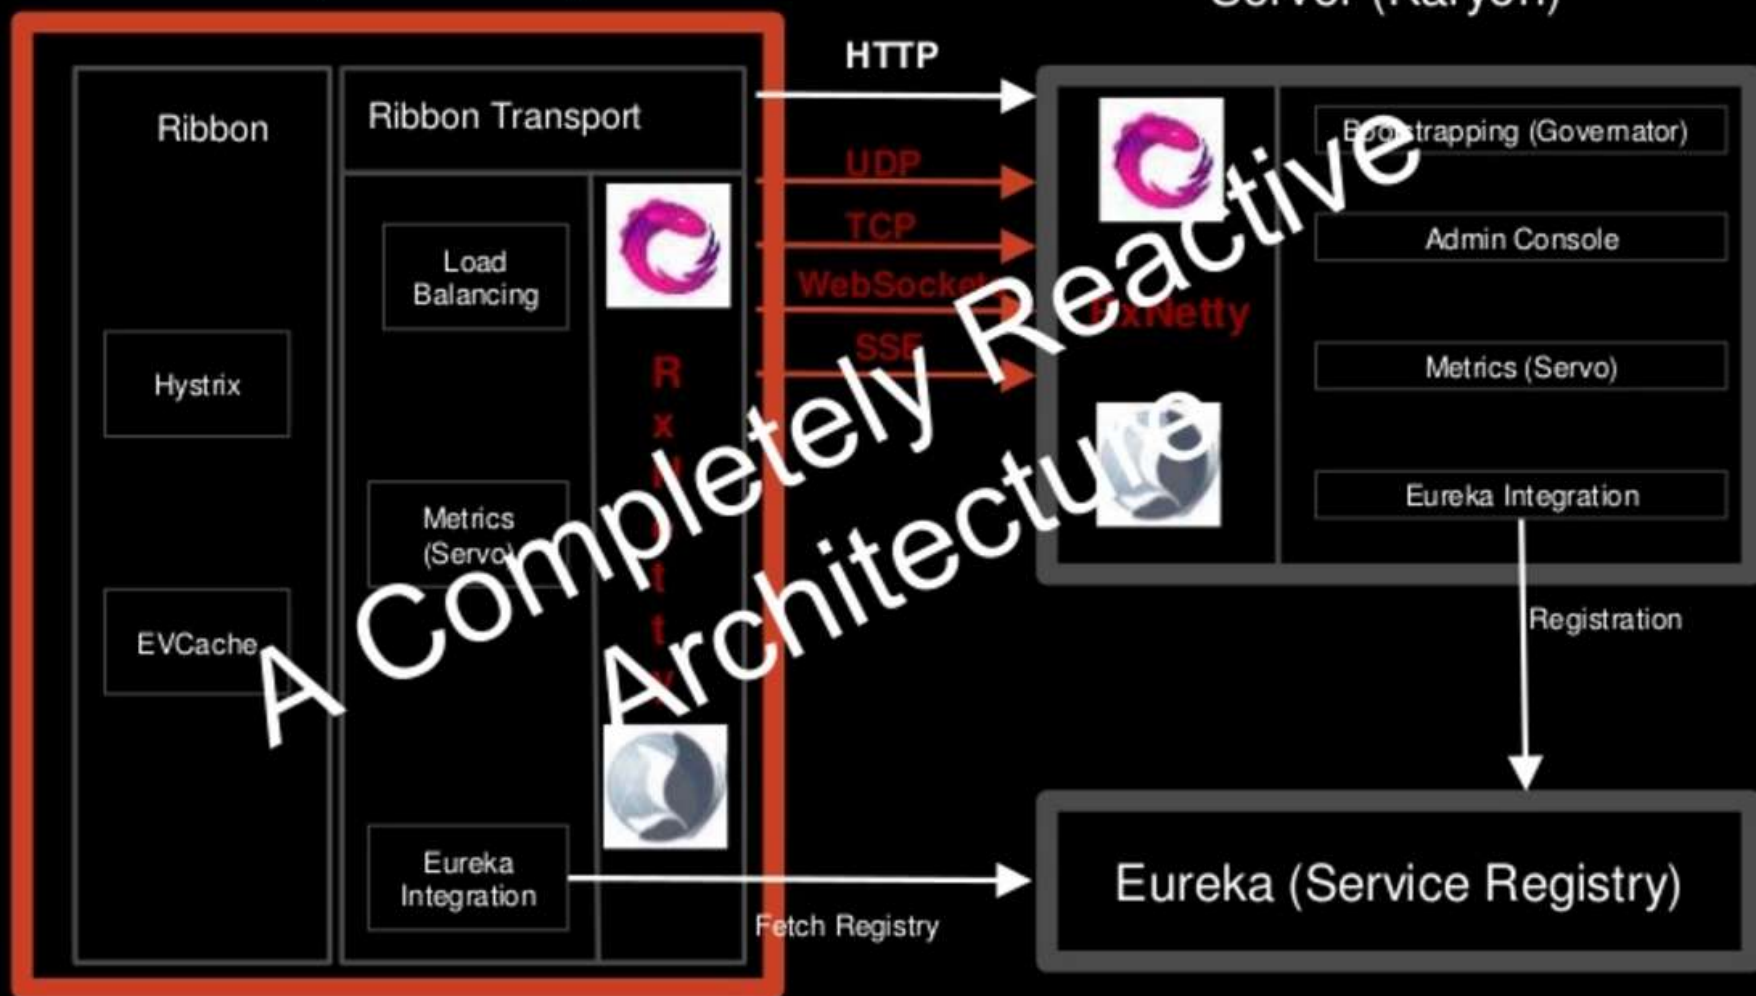
Server (Karyon)



Netflix IPC Stack (2.0)

Client (Ribbon 2.0)

Server (Karyon)



NETFLIX STACK

- Solid communications
 - Hystrix : latency and fault tolerance library
 - Eureka : registry for resilient mid-tier load balancing and failover
 - Ribbon : client based smart load balancer
 - Servo : monitoring library
 - EVCache : distributed in-memory data store for AWS EC2
 - RxNetty : reactive extension adaptor for netty
 - Karyon : blueprint of a cloud ready microservice
 - Zuul : edge gateway
 - Falcor : js library for efficient data fetching

NETFLIX STACK

- Automation
 - Asgard : specialized AWS console (app deployments, management)
 - Spinnaker : microservices console (clusters, pipelines), not opensourced yet
 - Atlas : near real-time operational insights
 - Vector : exposes hand picked high resolution metrics from PCP – Performance Co-Pilot hosts
 - SimianArmy : services (Monkeys) in the cloud for generating various kinds of failures, detecting abnormal conditions, and testing our ability to survive them
 - Dependencies automatically documented from from real traffic analysis

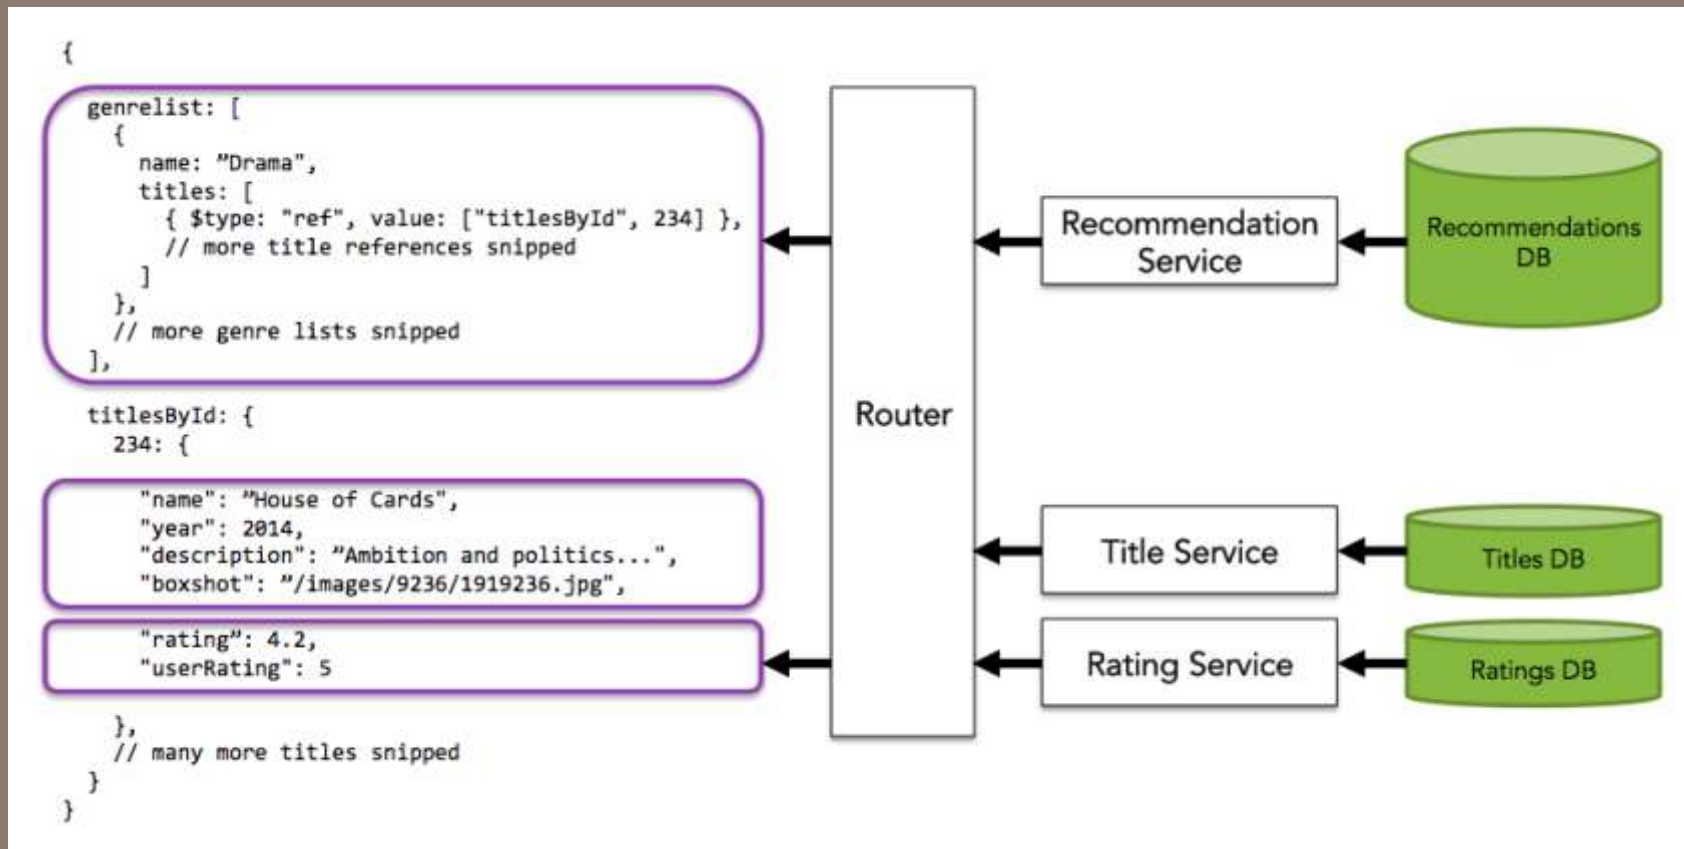


Netflix / eureka

- Mid-tier load balancing
 - a REST (Representational State Transfer) based service that is primarily used in the AWS cloud for locating services
 - for the purpose of load balancing and failover of middle-tier servers,
 - the load balancing happens at the instance/server/host level
- Comes with a Java-based client
 - The client also has a built-in load balancer
 - The client instances know all the information about which servers they need to talk to.
 - Does basic round-robin load balancing.
 - At Netflix, a much more sophisticated load balancer wraps Eureka to provide weighted load balancing based on several factors like traffic, resource usage, error conditions to provide superior resiliency.

NETFLIX FALCOR

- Working with JSON Virtual Resources



SPINNAKER - MICROSERVICES CONSOLE

SPINNAKER Applications Infrastructure Data

apiproxy-prod

Feedback

What's New
2015-08-18 08:13:38 PDT

apiproxy

PIPELINES

CLUSTERS

LOAD BALANCERS

SECURITY GROUPS

PROPERTIES

TASKS

CONFIG

Server Groups

Show ☒ Instances ☐ with details

Create Server Group

Filtered by: SEARCH: apiproxy-prod ACCOUNT: prod REGION: us-east-1 Clear All

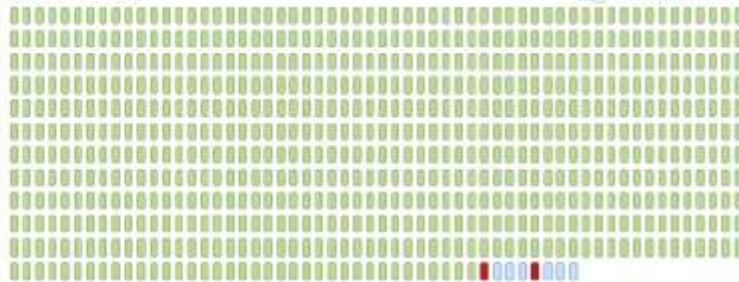
PROD apiproxy-prod

1357 ▲ / 9 ▼ / 15 ▲ / 639 ▼ : 98%

US-EAST-1

V136: Build: #1376

675 ▲ / 2 ▼ / 6 ▼ : 98%



V135: Build: #1291

250 —



Permalink

PROD apiproxy-prod-coalmine

1 ▲ : 100%

US-EAST-1

V936: Build: #1291

1 ▲ : 100%



Permalink

apiproxy-prod-v135

Server Group Actions

Insight

[SERVER GROUP IS DISABLED]

SERVER GROUP INFORMATION

Created: 2015-08-18 11:09:18 PDT

Id: prod us-east-1

VPC: None (EC2 Classic)

Availability Zones: us-east-1c
us-east-1d
us-east-1e

SIZE

Min: 250

Desired: 250

Max: 1359

Current: 250

Resize Server Group

View Scaling Activities

HEALTH

Instances: 250 —

LAUNCH CONFIGURATION

Name: apiproxy-prod-v135-20150818180918

Image ID: ami-bb9425d0

Image Name: 179727101194/edgezuul-1.27-h1291.1292fa0-x86_64-201508180743-trusty-hvm-sriov-efs

Base Image Name: trust-base-efs-64-201508080708-aws

SPINNAKER - MICROSERVICES CONSOLE

SPINNAKER Applications Infrastructure Data

Search

Feedback

What's New
(2015-09-15 08:21:26 PDT)

apiproxy

PIPELINES

CLUSTERS

LOAD BALANCERS

SECURITY GROUPS

PROPERTIES

TASKS

CONFIG

Group by Time Boundary

Pipeline Status

Not Started

Running

Completed

Failed

Canceled

Suspended

Show 5 per group

Configure Pipelines

Last Month

Trigger

Status

Enable Smoke Tested ASG

COMPLETED

9 minutes

Pipeline
DEV (cfieber testing)
[anonymous]
2015-06-25 10:54:15 PDT

WAIT FOR APPLICATION STARTUP

POST DEPLOY SMOKE TEST
(30 / 0)

ENABLE SERVER GROUP

DESTROY SERVER GROUP

Enable Smoke Tested ASG

COMPLETED

9 minutes

Pipeline
DEV (cfieber testing)
[anonymous]
2015-06-25 10:51:31 PDT

WAIT FOR APPLICATION STARTUP

POST DEPLOY SMOKE TEST
(36 / 0)

ENABLE SERVER GROUP

DESTROY SERVER GROUP

DEV (cfieber testing)

Manual Start

Build #1161

cfieber@netflix.com

2015-06-25 10:49:10 PDT

COMPLETED

14 minutes

MULTI-REGION ...

DEPLOY US-EAST-1

DEPLOY US-WES...

DEPLOY EU-WES...

SMOKE TEST BE...

SMOKE TEST BE...

SMOKE TEST BE...

DEV (cfieber testing)

Manual Start

Build #1161

cfieber@netflix.com

2015-06-25 10:39:23 PDT

FAILED

Smoke Test Before
Enable (us-east-1)

MULTI-REGION ...

DEPLOY US-EAST-1

DEPLOY US-WES...

DEPLOY EU-WES...

SMOKE TEST BE...

SMOKE TEST BE...

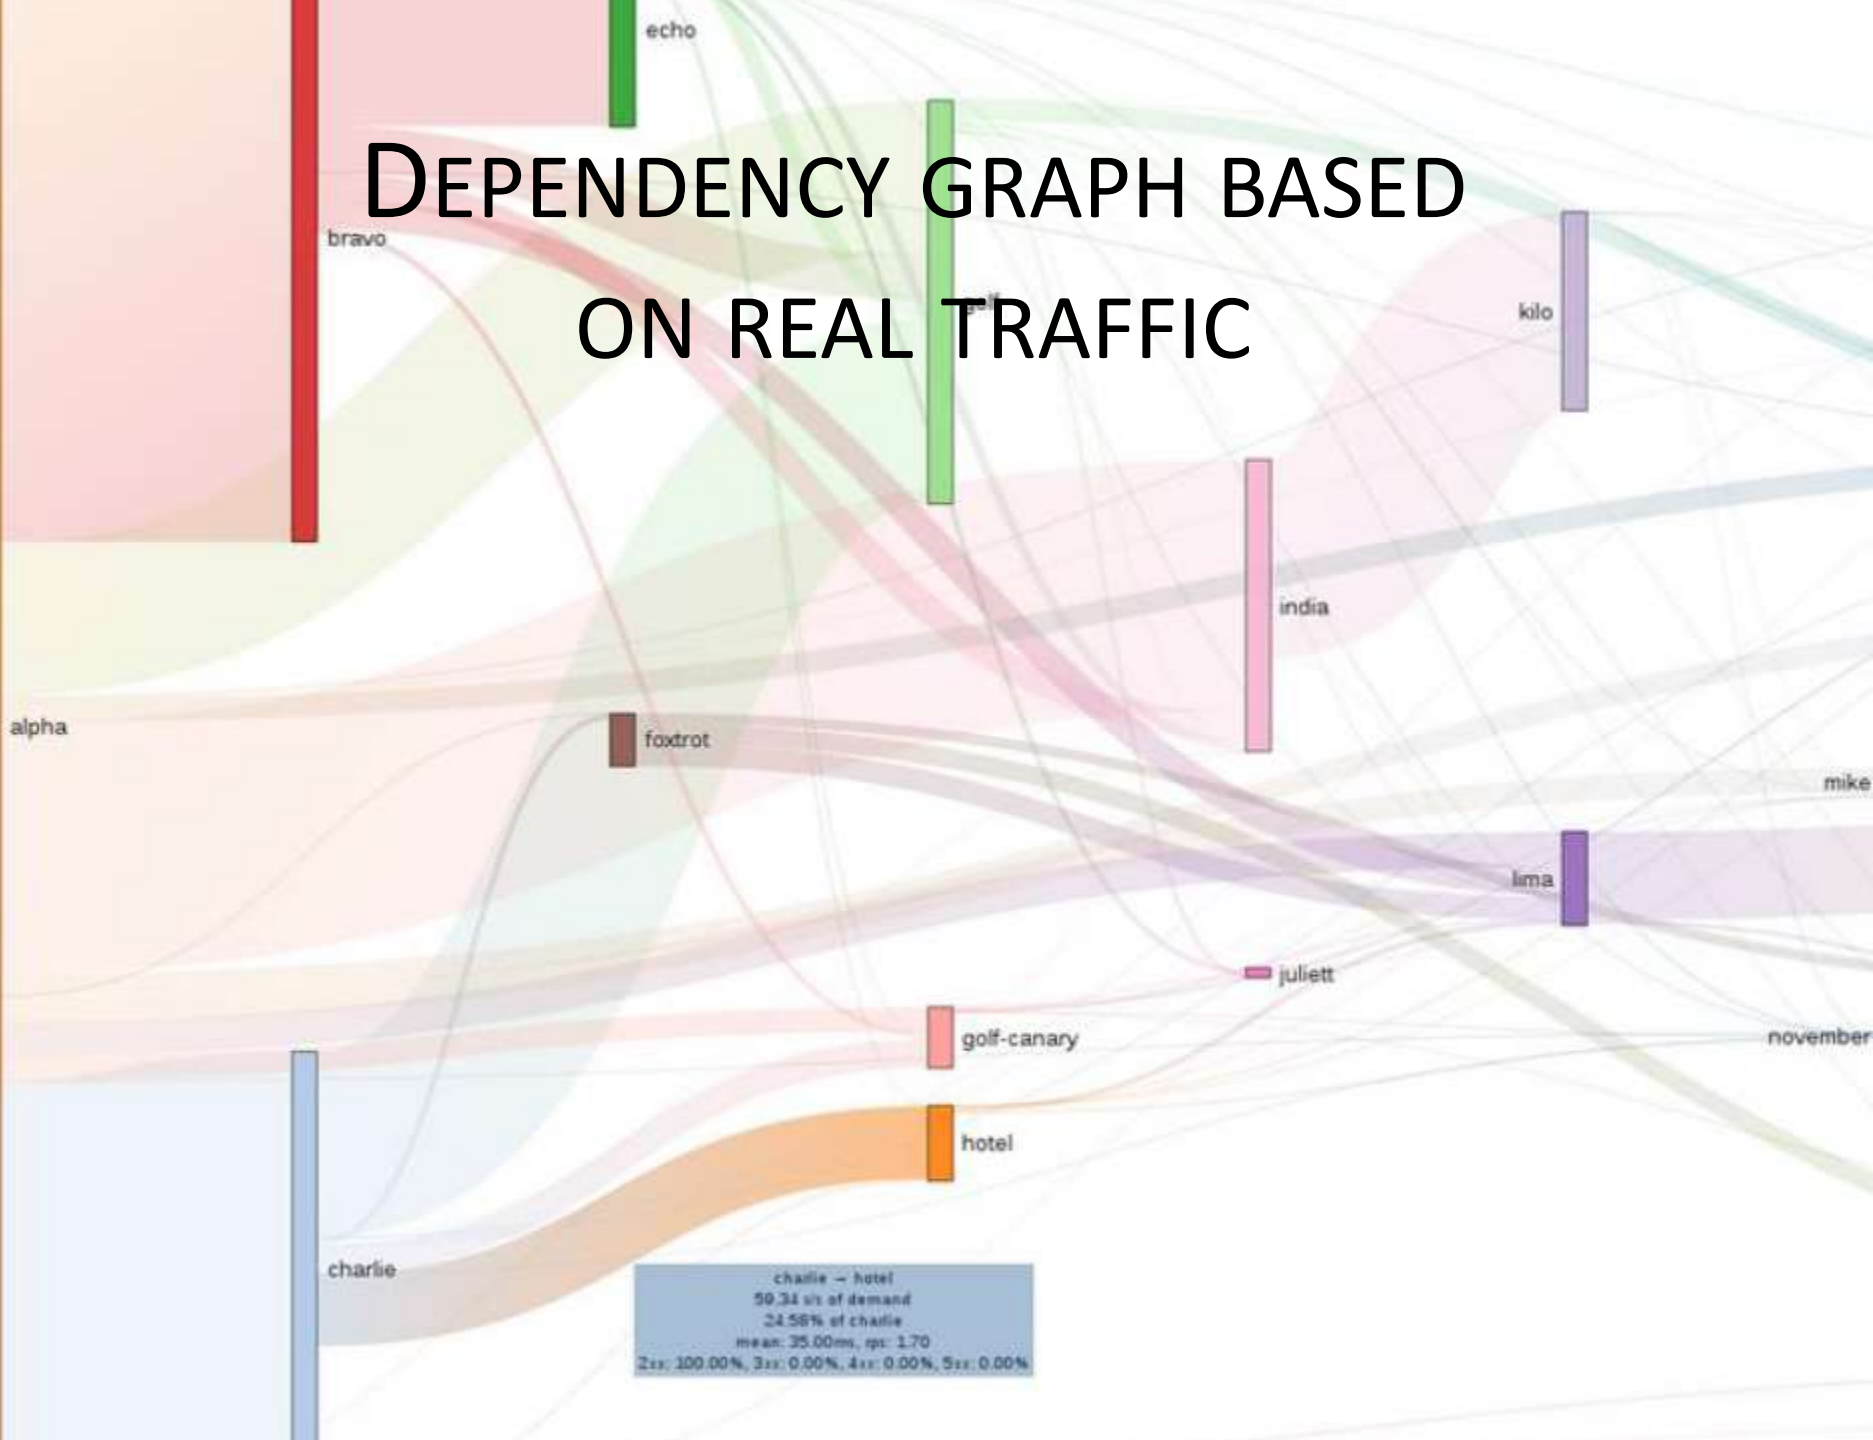
SMOKE TEST BE...

developer-michaels

Triggered Build

FAILED

DEPENDENCY GRAPH BASED ON REAL TRAFFIC



Graph

Y-Axis

Expression

Remove

Add...

Refresh

Graph

PNG

Query

Query:

cpuxawidle

Tags

Selecting an application, cluster, or ASG first will restrict other selections to meaningful values. At least one metric name is required for all queries.

nf.app / nf.cluster / nf.asg / nf.job:

api

- nf.app:api
- nf.app:apiconfig
- nf.app:apidaemon
- nf.app:apiproxy
- nf.app:atlas_alert_api
- nf.app:cass_api_multiregion
- nf.app:cas-api-us-east

name:

cpu

- CpuRawNice
- CpuRawSystem
- CpuRawUser
- CpuRawWait
- _ProcessCpuLoad
- _ProcessCpuTime
- SystemCpuLoad

Add tag:

Query settings

These settings modify the query results in a variety of ways.

Group by:

* name

Transform:

%

÷


x

+

-

>

<



CpuRawIdle

Max :	54.047	Min :	41.939
Avg :	50.182	Last :	44.451
Tot :	9.033k	Cnt :	180.000

CpuRawInterrupt

Max :	391.771μ	Min :	309.545μ
Avg :	351.041μ	Last :	332.947μ
Tot :	63.187m	Cnt :	180.000

CpuRawKernel

Max :	0.000	Min :	0.000
Avg :	0.000	Last :	0.000
Tot :	0.000	Cnt :	180.000

CpuRawNice

Max :	137.790m	Min :	457.534μ
Avg :	49.593m	Last :	78.929m
Tot :	8.927	Cnt :	180.000

CpuRawSystem

Max :	1.955	Min :	1.290
Avg :	1.417	Last :	1.498
Tot :	255.149	Cnt :	180.000

CpuRawUser

Max :	56.026	Min :	44.621
Avg :	48.297	Last :	53.952
Tot :	8.693k	Cnt :	180.000

CpuRawWait

Max :	324.259m	Min :	9.367m
-------	----------	-------	--------

Hostname ec2-54-159-49-252.compute-1.amazonaws.com



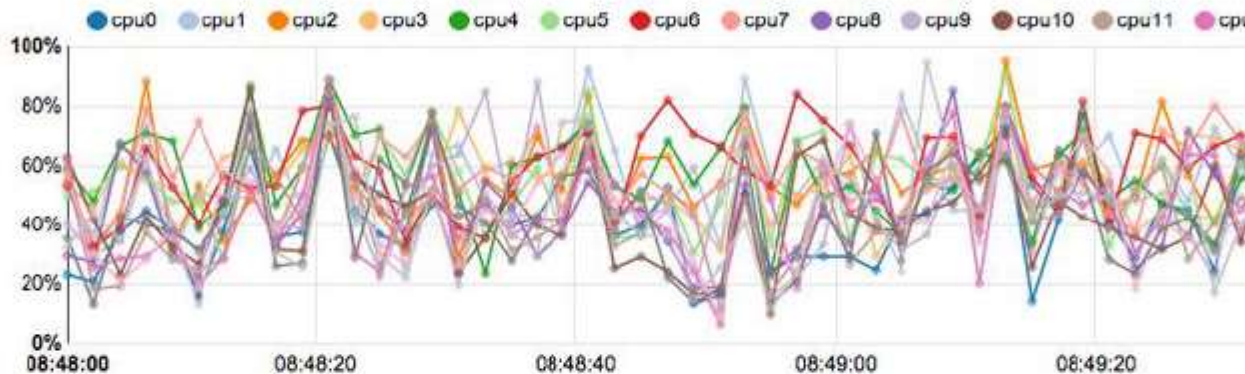
Widget ▾

Window 2 min

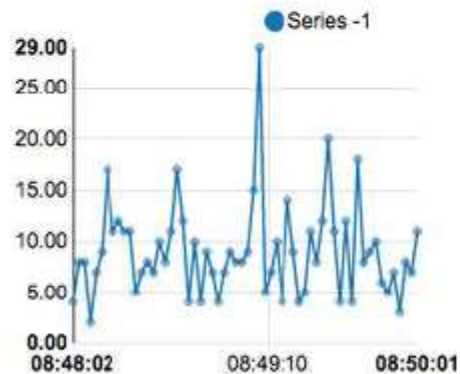
CPU Utilization



Per-CPU Utilization




Runnable



Load Average



Network kB



Network Packets



TCP Retransmits



Memory Utilization

How big should be my microservices ?

RESTish microservices ?

Isn't it SOA ?

The end of the monoliths ?

FAQ

HOW BIG ?

- Microservices ownership implies that each team is responsible for the entire lifecycle
 - functional vs divisional organizations
 - product management, development, QA, documentation, support
- Sizing depends on the system you're building
 - Amazon 2PT principle - Two Pizza Teams
 - 6 to 10 people to build/deploy/maintain a microservice
 - an average microservice at Gilt consists of 2000 lines of code, 5 source files, and is run on 3 instances in production

RESTish MICROSERVICES ?

- REST
 - Web lingua franca, 100% interoperable
 - development cost is generally higher
 - best practices : provide client sdks, (ex : generated from Swagger/RAML or other API description languages)
 - performance issues if not well-designed (chattiness)
 - best practices : experience based and coarser grained APIs
- RPC
 - optimized communication via binary formats
 - automated generation from IDL, polyglot by default
 - integrated support multiples scenarios : request/response, streaming, bi-directional streaming

HOW RESTISH MICROSERVICES ?

- RPC vs REST style depends on the system you're building and teams existing skills set
- Whatever the style, your microservices architecture **MUST** provide
 - Services Discovery,
 - Reliable Communications,
 - Operational insights (logs, monitoring, alerts, real time analysis)

GOT IT, BUT ISN'T IT SOA ?

- SOA so what ?
 - Enterprise SOA
 - Event-driven architecture (Pub/Sub)
 - Streaming Services (real-time time series, bidirectional)
 - Container-Services (ala Docker)
 - Nanoservices (ala AWS Lambda)
- Simply stated : Microservices are a SOA style for systems whose first goal is to scale
 - ⇒ in details, let's see how microservices differ from...

MICROSERVICES VS ENTERPRISE SOA

- Enterprise SOA is often seen as
 - multi-year initiatives, costs millions
 - complex protocols with productivity and interoperability challenges
 - central governance model that inhibits change
- Enterprise SOA is more about integrating siloed monoliths
 - generally via a smart and centralized service bus
- Microservices is scalable SOA
 - an architectural style to design, develop and deploy a large and complex system, so that it is easier to scale and evolve

VS ENTERPRISE SOA

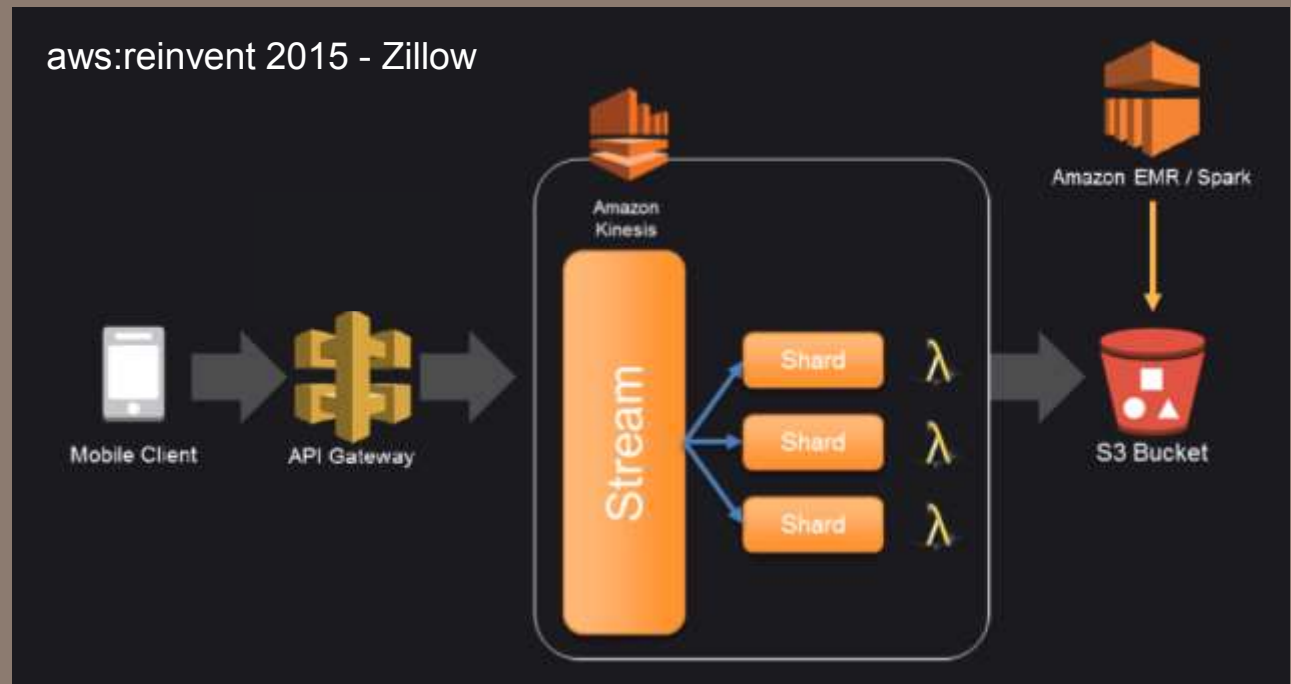
- ✓ Componentization via Services
- ✓ Organized around Business Capabilities
 - Products not Projects
 - Smart endpoints and dumb pipes
 - Decentralized governance
 - Decentralized data management
 - Infrastructure automation
- ✓ Design for failure
- ✓ Evolutionary design

VS EVENT DRIVEN ARCHITECTURE

- EDA fits well with Document oriented systems and information flows
- Communication between microservices can be a mix of RPC (ie, P2P calls) and EDA calls
- See EDA as a communication pattern for your microservices
- Can address choreography, orchestration, pipeline requirements

VS STREAMING SERVICES

- Streaming services fit well
 - if you have large volumes (log entries, IoT),
 - and/or if you aim at real time analysis
- Data ingestion endpoint of a microService
 - Real time analysis of a mobile app usage



VS CONTAINER SERVICES

- Containers provide the infrastructure to deploy your microservices independently one from another
- See Container Services as a building block of your global microservices architecture

VS NANOSERVICES

- Nanoservices are small pieces of code (functions)
 - Example : AWS Lambda, Auth0 Webtasks
- A microservice may leverage 1+ nanoservices

SUMP UP

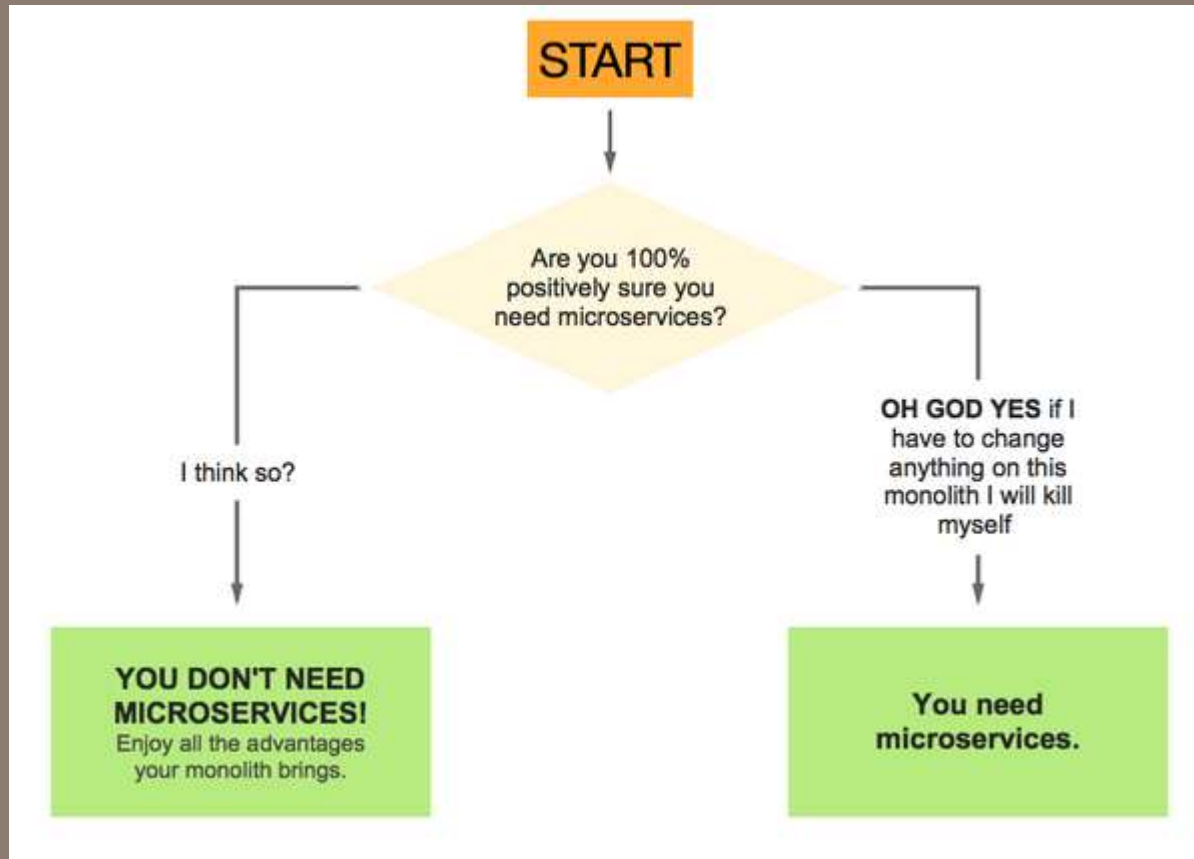
MICROSERVICES PATTERNS

- Solid communications
 - Fault tolerant libraries
 - Service discovery
- Committed teams
 - Devops culture
 - Code/Test/Deploy/Support, 24/7
 - Automation
- Ownership
 - Organisation aligned with the strategy
 - Insights via real time monitoring

MICROSERVICES IS A LONG JOURNEY

- Several years to implement
 - Communications, infrastructure, automation, monitoring, teams organization
 - The price to pay for extreme agility in a complex system at scale
- Prepare for iterative reworks
 - Multiples languages => maintainability
 - Numerous building blocks => updates & security
 - EOL of microservices stacks : manage your technical debt
 - Dependency Hell : keep control of your microservices segmentation
 - Ownership : is your organization ready ?

THE END OF MONOLITHS ?



FROM MONOLITHS TO MICROSERVICES

- Velocity of innovation for complex systems
 - Keep your monolith as is if you don't need to speed up features delivery
- To prepare for the journey
 - switch from layered architecture to internal APIs,
 - automate integration and deployment,
 - reorganize from divisional to functional teams committed to business and owning their code

REFERENCES

Microservices, Martin Fowler, James Lewis

- <http://martinfowler.com/articles/microservices.html>

A day in the life of a Netflix Engineer using 37% of the internet

- <http://fr.slideshare.net/AmazonWebServices/dvo203-the-life-of-a-netflix-engineer-using-37-of-the-internet>

Dzone – Getting started with Microservices

- <https://dzone.com/refcardz/getting-started-with-microservices>

SoundCloud – How we ended up with microservices

- http://philcalcado.com/2015/09/08/how_we_ended_up_with_microservices.html

gRPC - boilerplate to high-performance scalable APIs

- <http://fr.slideshare.net/AboutYouGmbH/robert-kubis-grpc-boilerplate-to-highperformance-scalable-apis-codetalks-2015>

To Go FURTHER

